

CME 213

SPRING 2019

Eric Darve

Homework 2

- Will be posted after class
- You can do most of the homework with your own computer, FarmShare, or the Google Compute Platform.
- See instructions in the homework assignment and on the class web page.
- Short demo of GCP

C++ threads, Pthreads

- Creating thread, thread join
- Mutex: locked/unlocked; used to protect access to shared variables; enforce order in read/write.
- Condition variables:
 - › Used to allow threads to become idle and wake up when a condition becomes true.
 - › `cond_wait`
 - › `cond_signal`

OpenMP

Pthreads/OpenMP

- Pthreads gives you maximum flexibility.
- It's a low level API that allows you to implement pretty much any parallel computation exactly the way you want it.
- However, in many cases, the user only wants to parallelize certain common situations:
 - › For loop: partition the loop into chunks and have each thread process one chunk.
 - › Hand-off a block of code (computation) to a separate thread
- This is where OpenMP is useful. It simplifies the programming significantly.
- In some cases, adding **one line** in a C code is sufficient to make it run in parallel.
- As a result, OpenMP is the standard approach in scientific computing for multicore processors.

OpenMP

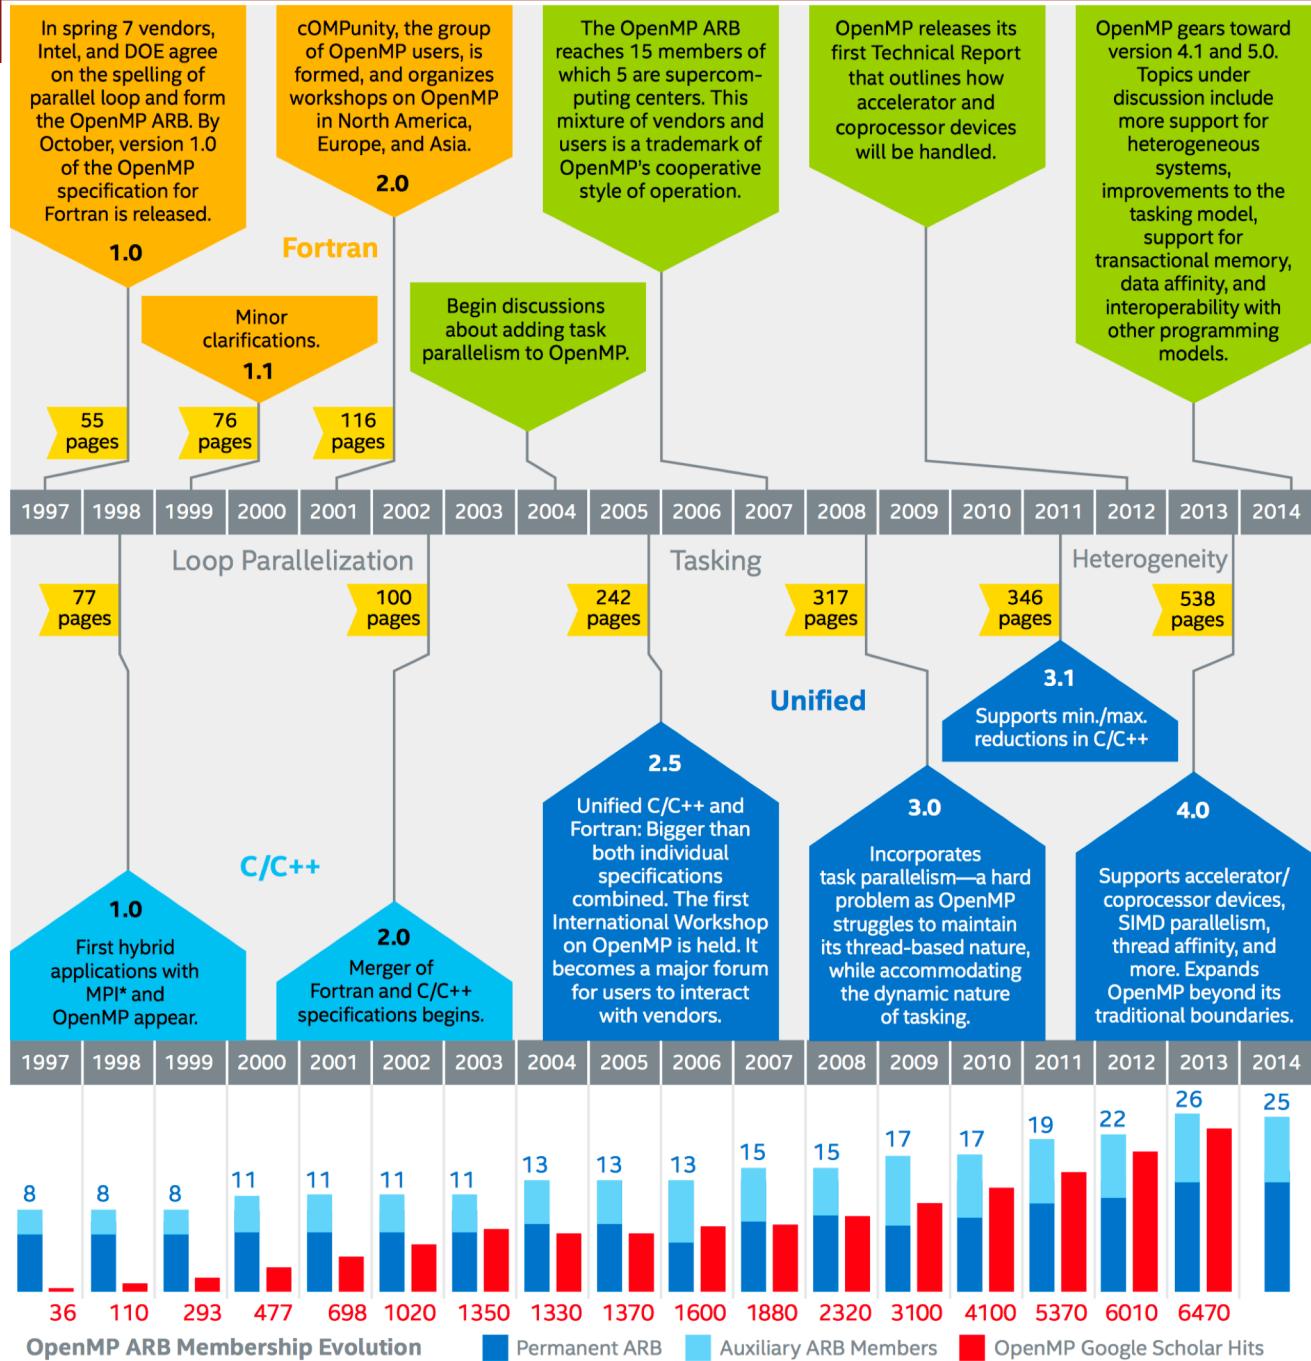
What is OpenMP?

- OpenMP is an Application Programming Interface (API), jointly defined by a group of major computer hardware and software vendors.
- OpenMP provides a portable, scalable model for developers of shared memory parallel applications.
- The API supports C/C++ and Fortran on a wide variety of architectures.

Hence, it is more portable and general than threads.

- OpenMP website: openmp.org
- Wikipedia: en.wikipedia.org/wiki/OpenMP
- LLNL tutorial <https://computing.llnl.gov/tutorials/openMP/>

- The Parallel Universe
- ARB: Architecture Review Board



Compiling your code

First things first

- Header file:

```
#include <omp.h>
```

- This is only needed if you explicitly use the OpenMP API.
- Compiler flags:

Compiler	Flag
icc	-openmp
icpc	
ifort	
gcc	-fopenmp
g++	
g77	
gfortran	

Parallel regions

For loop

For loops

- This is the most basic construct in OpenMP.
- You will need to use it for Homework Assignment 2.
- Let's go through an example together.
- Download the files from the class web page and open

for_loop.cpp

- I will show you a demo on GCP.

```
#pragma omp parallel for
    for (int i = 0; i < n; ++i)
    {
        x[i] = i;
        y[i] = 2 * i;
    }

print_vector(x);
print_vector(y);

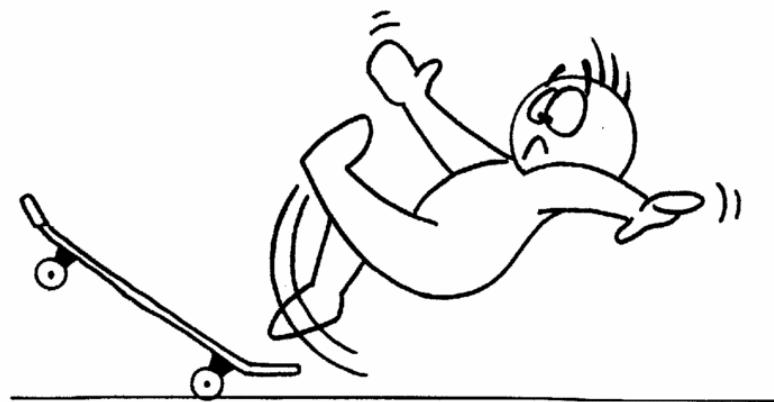
/* todo: write your own omp parallel for loop here */

print_vector(z);

for (int i = 0; i < n; ++i)
    assert(z[i] == x[i] + y[i]);
/* you must pass the test above if your for loop is correct. */
```

Directives

- OpenMP is based on directives.
- Powerful because of simple syntax.
- Dangerous because you may not understand exactly what the compiler is doing.

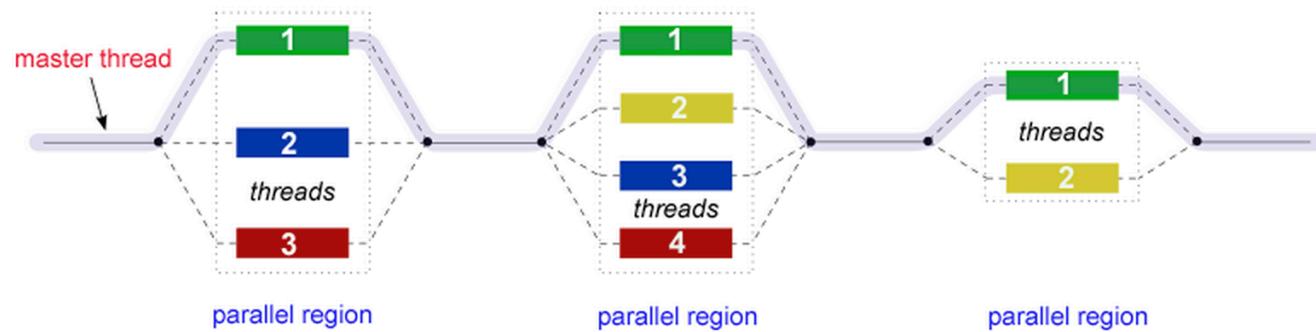


The most general concept: omp parallel region

The most basic directive is

```
#pragma omp parallel  
{ // structured block ... }
```

This starts a new parallel region. OpenMP follows a fork-join model:



Upon entering a region, if there are no further directives, a team of threads is created and all threads execute the code in the parallel region.

Basic example

`hello_world_openmp.cpp`

Compilation:

```
g++ -o hello_world hello_world_openmp.cc -fopenmp
```

```
#pragma omp parallel
{
    long tid = omp_get_thread_num();
    // Only thread 0 does this
    if (tid == 0)
    {
        int n_threads = omp_get_num_threads();
        printf("Number of threads = %d\n", n_threads);
    }
    // Print the thread ID
    printf("Hello World from thread = %ld\n", tid);
    // Compute some of the digits of pi
    DoWork(tid);
} // All threads join master thread and terminate
```

3. 1415926535897932384626433832795028841971693993751058209749445
923078164062862089986280348253421170679821480865132823066470938
446095505822317**253594081284811174502841027019385211055**59644
62294895493**038196442881097566593344612847564823378678**31652
712019091**45648566923460348610454326648213393607260249**14127
37245870**066063155881748815209209628292540917153643678**92590
3600113**3053054882046652138414695194151160943305727036**57595
919530**9218611738193261179**31051185480**7446237**996274956735188575
272489**1227938183011949129**83367336244**0656643**086021394946395224
737190**7021798609437027705**39217176293**1767523**8467481846766940513
2000**568127145263560827785**77134275778**9609173**6371787214684409012
24953430146549585371**05079227968925892354201**9956112129021960864
03441815981362977477**13099605187072113499999**9837297804995105973
1732816096318595024**459455**34690830264**2522308**2533446850352619311
8817101000313783875**2886587533208381420617177669147303598253490**
4287554687311595628**6388235378759375195778185778053217122680661**
300192787661119590**921642019893809525720106**54858632788659361533
818279682303019520**35301852968995773622599413891249721775283479**
13151557485724245**415069595082953311686172785588907509838175463**
7464939319255060**40092770167113900984882401**28583616035637076601
047101819429555**96198946767837449448255379774726847104047534646**
20804668425906**94912933136770289891521047521620569660240580381**
5019351125338**243003558764024749647326391419927260426992279678**
23547816360**09341721641219924586315030286182974555706749838505**
4945885869**269956909272107975093029553211653449872027559602364**
806654991**198818347977535663698074265425278625518184175746728**
90977772**7938000816470600161452491921732172147723501414419735**
68548161**3611573525521334757418494684385233239073941433345477**
6241686251**1898356948556209921922218427255025425688767179049460**
16534668049886272327917860857843838279679**766814541009538837863**
60950680064225125205173929848960841284886269456042419652850222
106611863067442786220391949450471237137869609563643719172874677

Pi algorithm

In our code, Pi is computed using:

$$\frac{\pi}{2} = 1 + \frac{1}{3} \left(1 + \frac{2}{5} \left(1 + \frac{3}{7} \left(1 + \frac{4}{9} \left(1 + \dots \right) \right) \right) \right)$$

Using this expansion, can you show that the code computes the digits of Pi,
4 at a time, assuming that:

$$\text{carry} + \text{sum} / \text{SCALE} < 10,000$$

- Is the previous algorithm parallel?
- Is this a good multicore implementation?
- How would you improve it?

- Computing pi in parallel is difficult.
- Many algorithms use sequential calculations using **high-precisions arithmetic**, that is you compute using numbers with a lot of digits.
- This leads to the natural question:

Is it possible to compute the
 n -th digit of π independently from the
others?

Bailey–Borwein–Plouffe formula

3. 1415926535897932384626433832795028841971693993751058209749445
923078164062862089986280348253421170679821480865132823066470938
44609550582231725359408128481117450284102701938521105559644
622948954930381964428810975665933446128475648233787831652
7120190914564856692346034861045432664821339360726024914127
3724587006606315588174881520920962829254091715364367892590
3600113305305488204665213841469519415116094330572703657595
9195309218611738193261179310511854807446237996274956735188575
2724891227938183011949129833673362440656643086021394946395224
73719070217986094370277053921717629317675238467481846766940513
20005681271452635608277857713427577896091736371787214684409012
24953430146549585371050792279689258923542019956112129021960864
03441815981362977477130996051870721134999999837297804995105973
17328160963185950244594553469083026425223082533446850352619311
88171010003137838752886587533208381420617177669147303598253490
42875546873115956286388235378759375195778185778053217122680661
30019278766111959092164201989380952572010654858632788659361533
81827968230301952035301852968995773622599413891249721775283479
1315155748572424541506959508295331686172785588907509838175463
74649393192550604009277016711390098488240128583616035637076601
04710181942955596198946767837449448255379774726847104047534646
2080466842590694912933136770289891521047521620569660240580381
5019351125338243003558764024749647326391419927260426992279678
2354781636009341721641219924586315030286182974555706749838505
4945885869269956909272107975093029553211653449872027559602364
80665499119881834797753566369807426542527862551814175746728
90977727938000816470600161452491927326391419927260426992279678
68548161361157352552133475741849468138523323907394133345477
624168625189835694855620992192218427356670179049460
165346680498862723279178608578438382796797766814541009538837863
60950680064225125205117392984960841284886269456042419652850222
106611863067442786220391949450471237137869609563643719172874677

$$\pi = \sum_{k=0}^{\infty} \frac{1}{16^k} \left(\frac{4}{8k+1} - \frac{2}{8k+4} - \frac{1}{8k+5} - \frac{1}{8k+6} \right)$$

Computing the n -th bit

This problem can now be reformulated as:
Can we compute the fractional part of

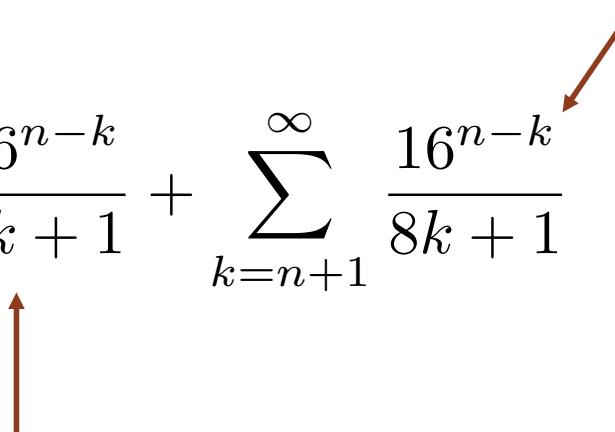
$$16^n \pi$$

Take:

$$\sum_{k=0}^{\infty} \frac{16^{n-k}}{8k+1} = \sum_{k=0}^n \frac{16^{n-k}}{8k+1} + \sum_{k=n+1}^{\infty} \frac{16^{n-k}}{8k+1}$$

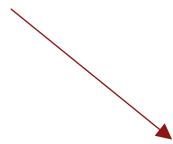
whole numbers
can be removed

Only a few terms are
needed



Computing the first sum

$$\sum_{k=0}^n \frac{16^{n-k}}{8k+1}$$



$$\sum_{k=0}^n \frac{16^{n-k} \mod(8k+1)}{8k+1}$$

This can be easily computed

Clause

- This is one of the tricky points of OpenMP.
- Recall in Pthreads that:
 - › Variables passed as argument to a thread are **shared** (they might be pointers in a **struct** for example)
 - › **Variables** inside the function that a thread is executing are **private** to the thread.
- OpenMP needs a similar mechanism: some variables are going to be shared (all threads can read and write), others need to be private.
- There are “complicated” rules to figure out whether a variable is private or shared.

Shared/private

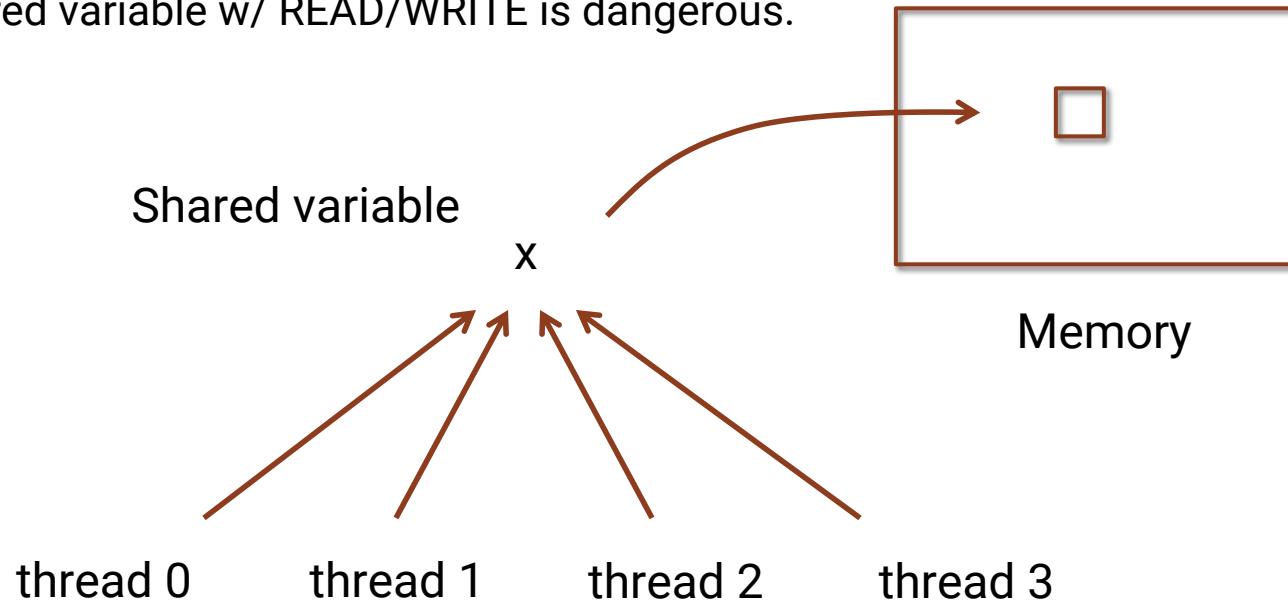
- See `shared_private_openmp.cpp`
- In a `parallel` construct, variables defined outside are shared by default.
- You can declare explicitly whether a variable is shared or private using

```
private(variable_name)
shared(variable_name)
```



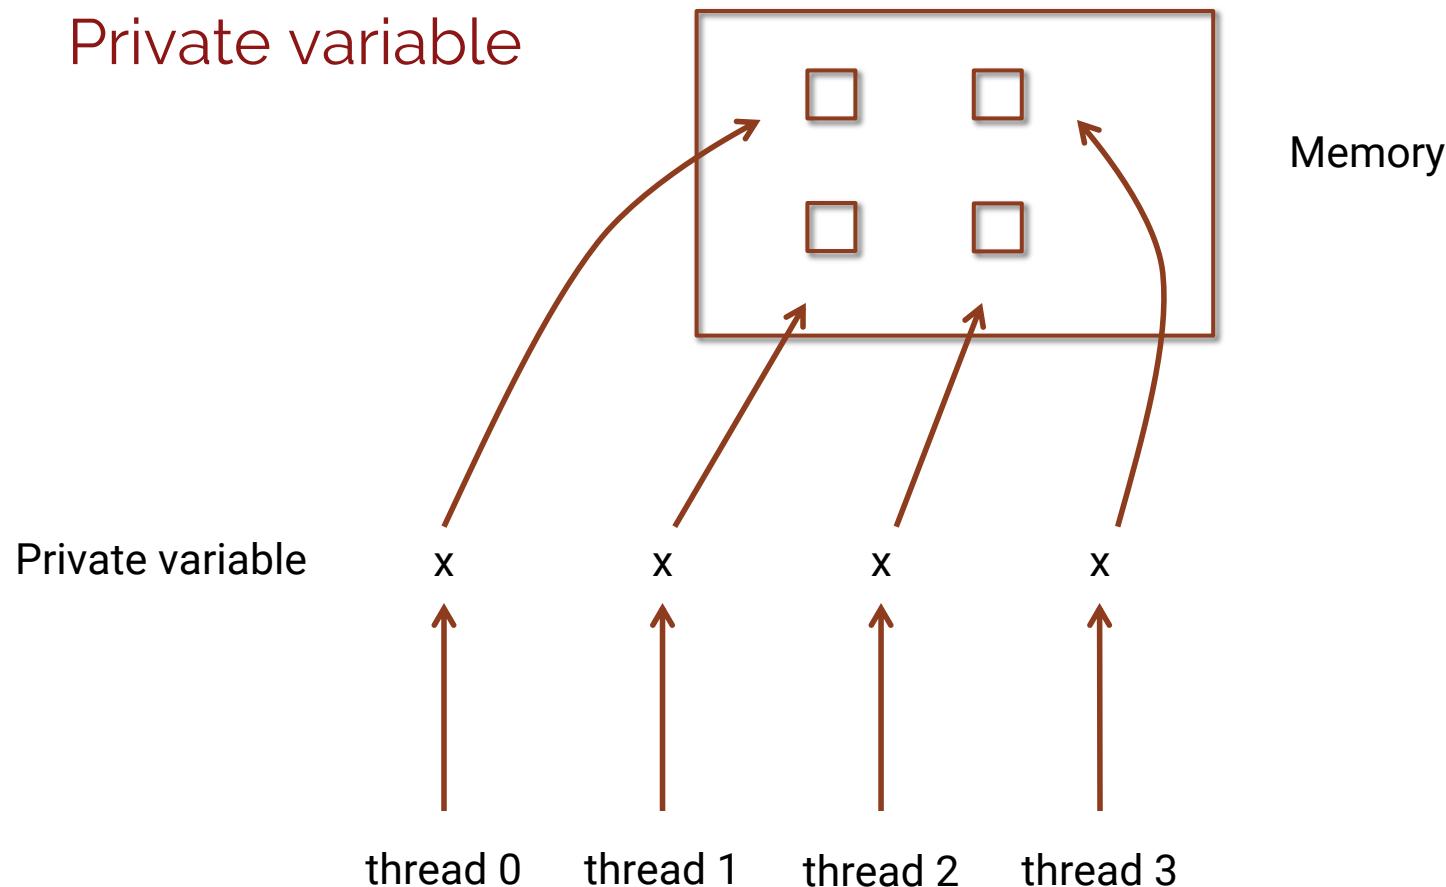
Shared variable

- Those are typically READ ONLY.
- Using a shared variable w/ READ/WRITE is dangerous.



Variable refers to the same memory location for all threads.

Private variable



Variable refers to a different memory location for each thread. Those variables are typically READ/WRITE.

Worksharing constructs

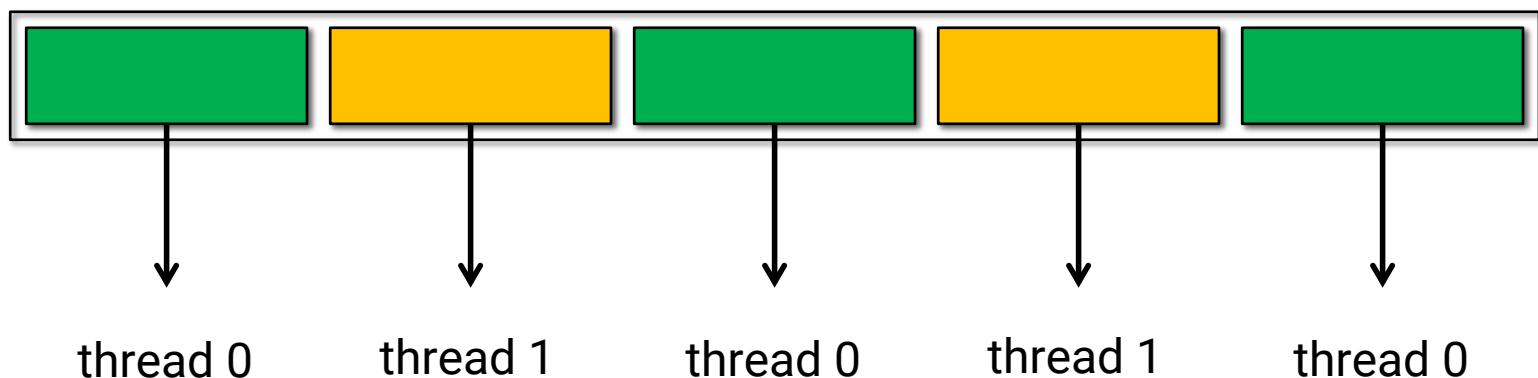


Parallel for loop

The most common approach to parallelize a computation on a multicore processor is to parallelize a `for` loop.

OpenMP has some special constructs to do that.

```
#pragma omp for [clause [clause] ... ]  
    for (i = lower bound; i op upper bound; incr expr) {  
    ...  
}
```



Example

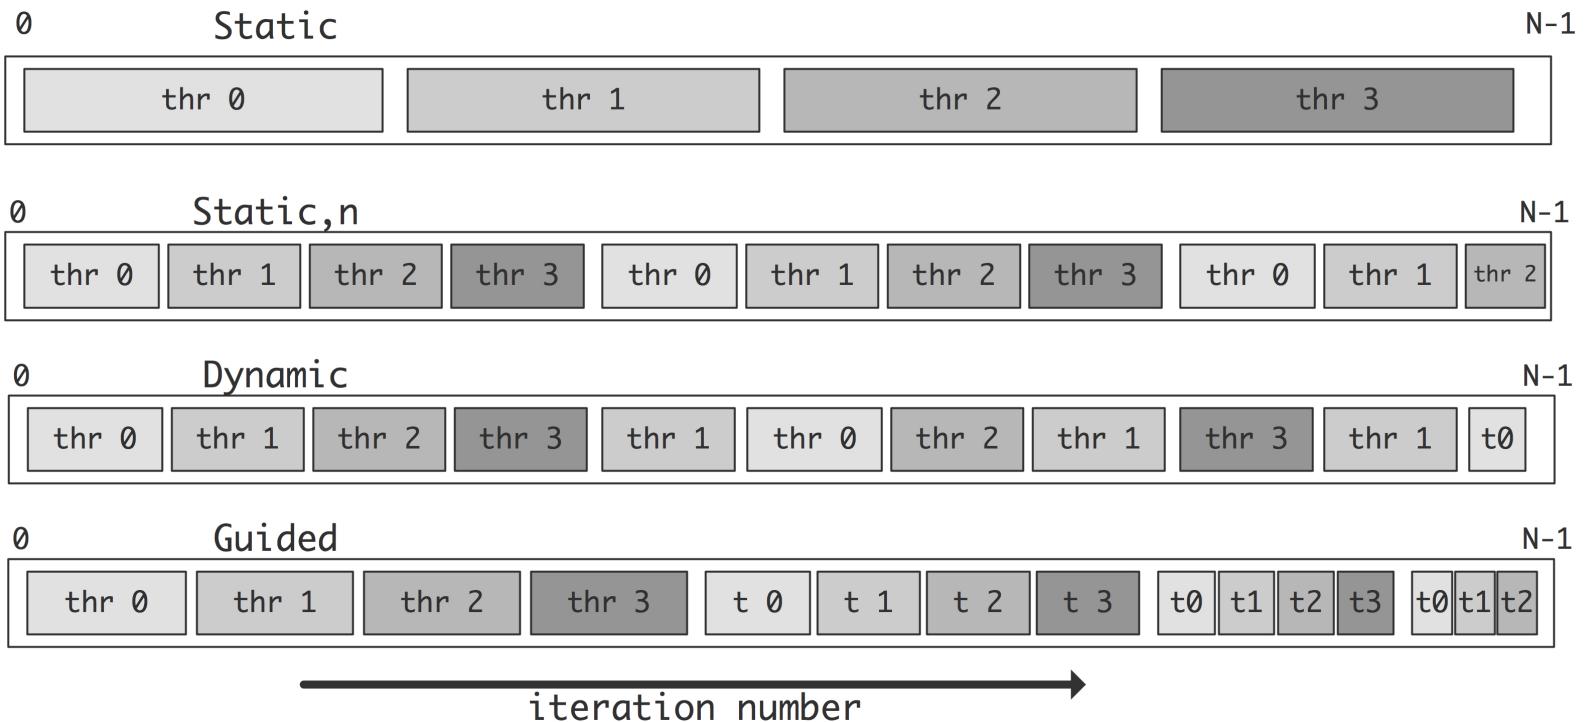
- Let's consider again the matrix-matrix example we used for Pthreads.
- See `matrix_prod_openmp.cpp`
- One line of code is sufficient to parallelize the calculation! This is the power of OpenMP.

```
$ ./matrix_prod_openmp -n 4000 -p 16  
$ top
```

```
#pragma omp parallel for
    for (int i = 0; i < size; ++i)
        for (int j = 0; j < size; ++j)
{
    float c_ij = 0;
    for (int k = 0; k < size; ++k)
{
    c_ij += MatA(i, k) * MatB(k, j);
}
mat_c[i * size + j] = c_ij;
}
```

Scheduling for loops

- How are the iterates in a for loop split among threads? This is important to fine-tune the optimization of your code.
 - This is a problem of load-balancing: how should we distribute the work so that we minimize the total execution time?
1. `schedule(static, block_size)`: iterations are divided into pieces of size `block_size` and then statically assigned to threads. This is the best **default** option.
 2. `schedule(dynamic, block_size)`: iterations are divided into pieces of size `block_size`, and dynamically scheduled among the threads; when a thread finishes one chunk, it is dynamically assigned another. This is useful when the **work per iteration is irregular**.
 3. `schedule(guided, block_size)`: specifies a dynamic scheduling of blocks but with decreasing size. It is appropriate for the case in which the **threads arrive at varying times** at a for construct (with each iteration requiring about the same amount of work).



Other worksharing constructs: sections

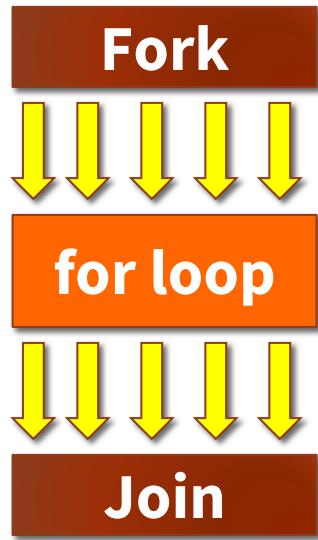


- There are situations where two independent pieces of work can be executed concurrently. For example, we may need to update two vectors independently.
- In that case, we would like to assign one thread to do each operation in parallel.
- This can be done using sections.
- The compiler is allowed to schedule the execution of the code inside each section concurrently.

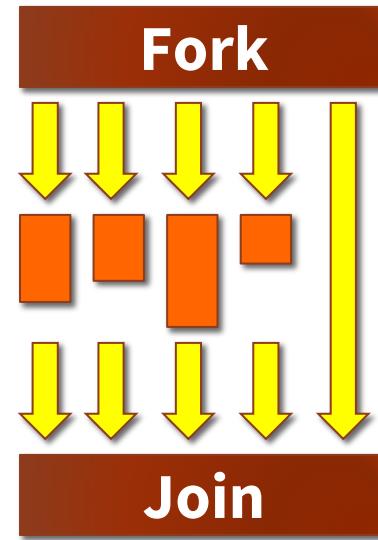
See `section.cpp`

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            for (unsigned i(0); i < size; i++)
                for (unsigned k(0); k < inner; k++)
                    c[i] += a[(i * k) % size];
        }
        #pragma omp section
        {
            for (unsigned i(0); i < size; i++)
                for (unsigned k(0); k < inner; k++)
                    d[i] += a[(b[i] + i * k) % size];
        }
    } // end of sections
} // end of parallel block
```

Summary



Large number of iterates.
Parallel for loop



Small and fixed number of independent tasks.
Parallel sections